

A Survey of Security Threats in Federated Learning

Yunhao Feng

National University of Defense Technology
Changsha, China

fengyunhaonudt@nudt.edu.cn

Jialiang Liu

National University of Defense Technology
Changsha, China

secondauthor@i2.org

Yinjian Hou

National University of Defense Technology
Changsha, China

secondauthor@i2.org

Yanming Guo

National University of Defense Technology
Changsha, China

secondauthor@i2.org

Abstract

Federated learning is a distributed machine learning paradigm that emerged as a solution to the need for privacy protection in artificial intelligence. Like traditional machine learning, federated learning is threatened by multiple attacks, backdoor attacks, Byzantine attacks, and adversarial attacks. The weaknesses are exacerbated by the inaccessibility of data in federated learning, which makes it more difficult to defend against these threats. This points to the need for further research into defensive approaches to make federated learning a real solution for distributed machine learning paradigm with securing data privacy. This paper aims to enhance understanding of the threats faced by federated learning and their defense mechanisms, and assist the academic and industrial communities in developing more robust federated learning systems. Our survey provides a taxonomy of these threats and defense methods, describing the general situation of this vulnerability in federated learning. We also sort out the relationship between these methods, their advantages and disadvantages, and discuss future research directions regarding the security issues of federated learning from multiple perspectives.

1. Introduction

Artificial intelligence (AI) can analyze large amounts of data and solve complex problems in various fields. Most existing state-of-the-art AI techniques rely on rich and high-quality datasets to train a highly accurate machine learning model. For example, while training the SAM model, Alexander et al. [40] build the largest segmented dataset ever, building more than 1 billion masks on 11 million images. However, in more practical reality, most companies cannot construct such large and highquality datasets as

Meta when training their model. To address this problem, one possible solution is the sharing of data among multiple organizations or companies to collaboratively train models.

But collaborative training presents several new challenges while it has developed rapidly. The first challenge stems from data privacy. The data held by organizations may contain sensitive information, such as medical or financial data of their users. Additionally, new legal frameworks are increasingly emphasizing the protection of individual data privacy [81]. Privacy protection introduces restrictions that prevent privacy data from leaving its originating organization and being uploaded to central server. The second challenge is related to the diverse organizations of data. Data from different organizations provides an increase in available data, driving advancements in artificial intelligence models. However, due to the various organizations, storing and processing data from different organizations, as well as reducing communication costs between these organizations have become a new challenge.

In response to such challenges, Google [54] develops a distributed machine learning framework called federated learning (FL). This framework allows each client to collectively train models without sharing their data. The data of each client remains private and inaccessible to others during training process. As shown in Fig. 1, in a typical federated learning framework, the server firstly sends a global model to all selected clients as local models. These clients then use their local datasets to train the local models and upload their trained model updates to the central server. After receiving the updates from all selected clients, the server updates the global model by averaging the uploaded updates. And the server sends the updated global model to the selected clients. Completing the four steps, one iteration of federated learning is complete. The entire federated learning requires multiple iterations to make model achieve con-




Figure 1. A schematic of federated learning. Each iteration of federated learning can be divided into four steps. 1: The central server sending the initialized global model to the client. 2: The clients then train locally and submit the local updates to the server. 3: The server performs the model aggregation. 4: the server sends the aggregated model to the clients.

vergence. Throughout the training process, only the data owners have access to their local data. This approach ensures the protection of data from unauthorized access by other clients or the central server, while also reducing communication costs between the clients and the server [95]. Due to the advantages of FL, FL has been widely applied to many fields in recent years [7, 14, 19, 48, 50, 87]. However, further research has shown that federated learning also faces numerous security risks [21, 33, 69, 78, 103] like backdoor attacks, adversarial attacks, and Byzantine attacks.

We primarily enumerate backdoor attacks, adversarial attacks, and Byzantine attacks in federated learning, along with their corresponding defense methods [11, 18, 43, 51, 80, 98]. These attacks pose significant threats to federated learning systems, and in response to these new threats, several FL defense methods have been proposed. In previous surveys [21, 33, 69, 78, 103], researchers have explored some attacks and defense strategies in the context of federated learning. We do not believe that this classification between the two types of attacks can be perfect cause some attack methods [5, 76, 105, 107] start from the aspect of modifying the client model. These surveys often categorize backdoor attacks and Byzantine attacks as data poisoning attacks and rarely delve into adversarial attacks in the field of federated learning. In this survey, we propose a new classification approach for these two types of attacks and provide an overview of their corresponding defense mechanisms. We also supplement this paper with content on adversarial attacks and their defense mechanisms. Furthermore, we conduct a detailed analysis of the effec-




Figure 2. As shown in the figure, a common backdoor attack method is to insert backdoor samples during model training, so that the model can show the predicted results desired by the attacker on the backdoor samples during prediction.




Figure 3. In Byzantine attacks, there exist one or more malicious participants(client k) in the federated learning system who disrupt the training process by sending incorrect or misleading updates to the central server, causing abnormal convergence.

tiveness and limitations of these attack methods and their corresponding defense mechanisms. In section 2, we give a preliminary introduction to the definitions of the three attacks discussed in this survey and try to find their commonalities. In section 3-8, we introduce the current state of attack and defense methods for each of the three threats. Finally, we analyze the advanced research and problems faced by federated learning in section 9.

Table 1. Comparison of Attacks

Attack Category	Goal	Mechanism	Attack Phase	Targeted Attack
Backdoor Attack	• Present results as attackers expect on the backdoor samples. • Behave normally on benign samples.	Excessive learning ability of models.	Training	Targeted
Byzantine Attack	• Reduce model generalization. • Make model difficult to converge.	Distribution of federated learning clients.	Training	Untargeted
Adversarial Attack	• Misclassify attacked samples. • Behave normally on benign samples.	The difference of samples in feature space.	Inference	Targeted / Untargeted

2. Threat Models

In this section, we give a preliminary introduction to the definitions of the three attacks discussed in this survey.

Backdoor attacks [5, 28, 62, 77, 84] refer to a malicious backdoor added to the global model by malicious participants during training process. The backdoor can be triggered by specific inputs, allowing attackers to control the outputs of the model. The goal of backdoor attack is to make the model maintain correct outputs on benign samples, while presenting the bad results as attackers expect on backdoor samples as shown in Fig. 2, .




Figure 4. Adversarial attacks require adding subtle and carefully crafted perturbations to the input data to deceive the model and cause it to make incorrect predictions.

Byzantine attacks [22, 32, 66] is that there exist one or more malicious participants in the federated learning system who disrupt the training process by sending incorrect or misleading updates to the central server, causing abnormal convergence, as shown in Fig.3.

Adversarial attacks [13, 44, 101, 112] require adding subtle and carefully crafted perturbations to the input data to deceive the model and cause it to make incorrect predictions as shown in Fig.4.

These attacks have different goals, mechanisms, and attack stages. We can categorize these threats into two main

stages: the training phase and the inference phase. Additionally, we can also categorize them between untargeted attacks and targeted attacks based on whether a specific target is present or not. And we make a comparison among attacks in Tab. 1.

2.1. Attack in Training Phase & Attack in Inference Phase

(1)Attack in Training Phase: Attacks during the model training process are intended to impact the federated learning model. During the training stage, backdoor attacks involve the insertion of a backdoor into the model, whereas input deception models with triggers are utilized during the reasoning stage to cause the model to generate incorrect results [56, 99]. Byzantine attacks disrupt the convergence of the model by utilizing malicious clients or servers [22]. It causes the global model fail to converge by sending harmful information or disrupt communication.

(2)Attack in Inference Phase: Attacks that occur during the reasoning phase are typically intended to alter the model’s reasoning outcomes and deceive it into generating incorrect outputs [6]. Adversarial attacks leverage the model’s vulnerability to disturbances and utilize samples with adversarial perturbations as input to the model, causing it to produce erroneous outcomes.

2.2. Untargeted Attack & Targeted Attack

(1)Untargeted attacks: Untargeted attacks are designed to compromise the integrity of the target model in arbitrary manner. Byzantine attack is one form of untargeted attacks that involves uploading malicious gradients to the server in an arbitrary manner, with the goal of causing the global model to fail [8, 17, 41, 91].

(2)Targeted Attacks: A targeted attack is executed with the aim of inducing the model to produce the target label specified by the adversary for specific testing examples, while keeping the testing error for other testing examples unaffected [17]. Backdoor attack is a typical application of targeted attacks.

3. Backdoor Attack

Backdoor attacks on deep neural networks require malicious backdoors to be secretly implanted in the model. This

Table 2. Comparison of Backdoor Attacks

Attack Method	Category	Persistence	Mechanism	Stealthiness
Semantic Backdoor	Model Poisoning	Bad	Model replacement.	Bad
Label Flipping	Data Poisoning	Bad	Generated poisoned examples.	Bad
CLean Label	Data Poisoning	Outstanding	Sample camouflage.	Outstanding
Transferable Clean Label	Data Poisoning	Outstanding	Sample camouflage.	Outstanding
Backdoor with Edge Data	Data Poisoning	Normal	Long-tail distribution.	Normal
Distributed Backdoor	Data Poisoning	Bad	Embedded backdoors distributed from multiple clients.	Outstanding
Neurotoxin	Model Poisoning	Outstanding	Attack parameters that change less during training.	Normal
Optimization-based Backdoor Attack	Model Poisoning	Normal	Optimization method	Normal
Distance Awareness Attack	Model Poisoning	Normal	Attacking Distance-aware Attack	Normal
Dynamic Backdoor Attack	Data Poisoning	Normal	Dynamic backdoor	Outstanding
Chameleon	Data Poisoning	Outstanding	Adaptive technique.	Normal

enables the model to function normally when processing benign inputs, but triggers pre-defined malicious behavior when presented with a specific malicious trigger. The first neural backdoor in centralized settings can be traced back to 2013 [26, 30]. Due to the unique nature of federated learning, whereby the model is trained on individual clients, it is more susceptible to backdoor attacks compared to the general centralized training model. These backdoor attacks in federated learning can be divided into two categories based on the different stages at which the adversary inserts backdoors into the training pipeline: data poisoning attacks and model poisoning attacks.

3.1. Data Poisoning Attack

In data poisoning attacks, it is assumed that the adversary has full control over the training data collection process of compromised clients. Then the poisoned dataset typically consists of a combination of clean data with ground-truth labels and data with backdoor triggers that have targeted labels.

(1)Visible Poisoning: Early methods [5, 28] of backdoor attacks in federated learning can rely on a single trigger, meaning that all corrupted clients inject the same trigger with their local training datasets. The triggers are usually predetermined, such as a square located at redundant

pixels in the image. During reasoning, the inserted triggers are used on malicious clients to activate the aggregation model [5, 28]. While the effectiveness of inserting the backdoor has been shown to be significant, the aforementioned approach merely transfers backdoor attacks from centralized learning directly to federated learning, without fully leveraging the distributed nature of the latter. This is because the same triggers are embedded in all adversarial clients. Some studies [90] take advantages of federated learning to carry out backdoor attacks. As shown in Fig. 5, Xie et al. [90] propose a distributed backdoor attack (DBA) that decomposes a global trigger pattern, similar to a centralized attack, into local patterns and embeds them into different malicious clients. Compared to traditional methods that insert the same trigger, DBA is more efficient and covert due to its hidden local trigger mode, making it easier to bypass robust aggregation rules.

But distributed backdoors are also easy to detect because the triggers are static. And if the static triggers are changed to dynamic, the backdoors will be more difficult to detect. Salem et al. [70] conduct a clear and systematic study on the feasibility of dynamic flip-flops, which can facilitate the backdoor attacks by generating antagonistic network algorithms to create triggers. This way, the same tags can be hijacked with different trigger patterns that share similar




Figure 5. DBA decomposes a global trigger pattern, similar to a trigger in centralized attack, into local patterns and embeds them into different malicious clients.

potential representations and positions. Li et al. [45] also conduct a direct investigation of dynamic triggers. These triggers can be flexibly produced during the attack phase, as they maintain their effectiveness even under substantial changes. The results of the study show that dynamically triggered backdoor attacks are more powerful, and they require new techniques to be defeated because they break the static trigger hypothesis of most current defense systems. Despite the potential threats of data poisoning attacks in federated learning, they face many practical limitations due to the unique distributed characteristics [84]. This is because the data distribution and model aggregation steps in federated learning tend to neutralize most of the contributions of the backdoor model, leading to rapid forgetting of the backdoor by the global model. In light of this situation, Wang et al. [84] propose selecting poisoning samples from edge data to reduce the forgetting effect caused by model updating.

Dai et al. [16] propose a new backdoor attack method called Chameleon, which enables attackers to create more persistent visual backdoors by adapting to peer-to-peer images. The durability of the backdoor largely depends on the existence of two types of identical benign images that are closely related to the toxic image: Interferers, which are images that share the same original tag as the toxic image. And facilitators, which are images with the target back tag. Interferers can cause update conflicts between toxic updates and benign updates, which may reduce the accuracy of the backdoor. Conversely, facilitators can help reintroduce backdoor information into the federated learning model and mitigate the catastrophic forgetting effect after the attacker leaves the training process. Inspired by these observations, Chameleon is designed to amplify these effects and enhance

the durability of the backdoor.

(2) Invisible Poisoning: In this context, the term 'invisible' denotes the user's ability to execute a backdoor attack on a sample without requiring any additional actions to be performed on the sample itself. Label flipping is a widely recognized attack in centralized machine learning (ML), as demonstrated in previous research studies [73, 74]. In addition, it is also a suitable method for the federated learning (FL) scenario, given its adversarial goal and capabilities [79]. As shown in Fig. 6, some of the samples labeled 'dog' are flipped to 'cat' and the samples of 'cat' are flipped to 'dog'.




Figure 6. Label Flipping. In this picture, some of the samples labeled 'dog' are flipped to 'cat' and the samples of 'cat' are flipped to 'dog'.

Nevertheless, the label flipping technique has its limitations since it necessitates modifying the label of the samples, making it less practical. Thus, an attack technique that is more covert and can deceive manual inspection would be more appealing in this scenario. As shown in Fig. 7, clean-label attack preserves the label of the poisoned data, and the manipulated image still appears to be a benign sample [71, 110]. This type of attack leverages feature collision, where the crafted poison examples continue to resemble the source instances in visual appearance, while being closer to the targeted instances in the latent feature space. Data poisoning faces the risks of being easily detected and traced. At the same time, the initial assumption that the attacker has complete control over the data set is difficult to achieve in many scenarios.

3.2. Model Poisoning Attack

To address the limitations of data poisoning attacks, we can not only focus on improving the data poisoning technology itself, but also exploring the potential of model poisoning techniques. Since the average method is the most widely used approach for aggregating local updates from clients, a simple way to amplify the backdoor effect is to prioritize updates from adversarial clients over those from benign clients.

Bagdasaryan et al. [5] propose the first backdoor attack against federated learning. Their approach involves training




Figure 7. Clean-label attack preserves the label of the poisoned data, and the manipulated image still appears to be a benign sample. As you see, it's almost impossible to tell the difference between poisoned examples and benign examples.

a backdoor model that closely resembles the global model, which is then used to replace the latest global model. To improve the effectiveness of this replacement, they slow down the learning rate to extend the lifespan of the backdoor model, and add an anomaly detection term to the loss function to avoid detection. This strategy requires careful evaluation of the global parameters and performs better when the global model is close to convergence.

But similar to data poisoning, such direct substitutions are easily detected. In order to avoid such substitutions, Zhou et al. [105] propose an optimization-based model poisoning attack that involves injecting adversarial neurons into the redundant space of a neural network to maintain the attack’s concealment and persistence. To identify the redundant space, the Hessian matrix is used to measure the updated distance and direction of each neuron’s main task. An additional term is then added to the loss function to prevent poisoned neurons from being injected in locations that are particularly relevant to the main task. In a similar vein, Zhang et al. [107] propose a persistent backdoor attack called Neurotoxin. This method relies on the empirical observation that the norm of a stochastic gradient is primarily concentrated in a small number of ‘heavy hitter’ coordinates. Neurotoxin identifies these heavy hitters using the top-k heuristic and avoids them. By avoiding directions that are most likely to receive large updates from benign devices, the chance of the backdoor being erased is mitigated.

Sun et al. [76] propose a distance-aware attack (ADA), which enhances poisoning attacks by identifying optimized target classes in the feature space. They address the challenge of limited prior knowledge of customer data that competitors may face. To overcome this problem, ADA infers

the pairwise distances between different categories in the potential feature space from the shared model parameters using backward error analysis. They conduct an extensive empirical evaluation of ADA by varying attack frequency in three different image classification tasks. As a result, ADA successfully improves the attack performance by 1.8 times in the most challenging cases with attack frequency of 0.01x.

3.3. Summary Of Federated Backdoor Attack

As shown in Tab. 2, we make a comparison of backdoor attack methods mentioned in this section. An effective attack method should satisfy the three conditions of accuracy, persistence and stealthiness. These existing attack methods are improved in the following directions: using multiple malicious clients to insert backdoor attacks, combining generation technology with backdoor attacks to insert dynamic backdoors, and conducting research on persistent backdoor attacks in federated learning.

4. Defenses against Backdoor Attack

To mitigate the problem of backdoor attacks in federated learning, various defensive techniques have been proposed. As shown in Fig. 8, given that we previously categorized backdoor attacks as data poisoning attacks and model poisoning attacks, we will now discuss defensive strategies for each of these attack types.

4.1. Defense against Data Poisoning

The simplest approach is to filter out poisoned data samples, which aims to remove the poisoned samples from the training datasets. After the filtering process, only benign samples or purified poisoned samples are used during the training process, thus eliminating the creation of backdoors from the source.

Tran et al. [80] are the first to investigate methods for filtering out malicious samples from the training sets. They demonstrate that poisoned samples tend to leave detectable traces within the covariance range of the feature representation. Exploiting this insight, it is possible to filter out poisoned samples from the training sets. Chen et al. [11] propose a two-stage filtering approach. As shown in Fig. 9, in the first stage, activation values of samples in each class are clustered into two groups. And in the second stage, it is determined which clusters correspond to poisoned samples. This method is the first methodology for detecting poisonous data maliciously inserted into the training set to generate backdoors that does not require verified and trusted data. It has been released as a part of the IBM Adversarial Robustness ToolBox. Similarly, Zeng et al. [98] reveal that poisoned samples of existing attacks have some high-frequency artifacts even if their trigger patterns are invisible in the input space. Based on this observation, they design




Figure 8. This graph is a summary of some of the defensive solutions mentioned in this article, we have summarized and categorized some of them in this graph, and sorted them in chronological order.

a simple yet effective filtering method based on those artifacts. Based on data-driven defense methods, in addition to filtering out samples, it is also possible to consider directly preprocessing the samples, specifically by erasing any backdoors within them to prevent them from being embedded in the model. Liu et al. [51] propose a pretrained autoencoder as a preprocessor to prevent malicious samples from embedding backdoors by preprocessing the input samples without affecting data classification accuracy.

Doan et al. [18] propose a two-stage image processing method called Februus, in which, in the first stage, Februus uses GradCAM to identify influential regions, which are then removed and replaced with neutral color frames. Subsequently, as shown in Fig. 10, Februus uses a GAN-based repair method to reconstruct the masked regions to mitigate their adverse effects (such as benign accuracy reduction). Li et al. [47] discuss the properties of existing poisoning-based static trigger mode attacks. They demonstrate that the attack performance may sharply decrease if the appearance or location of the trigger is slightly changed. Based on this observation, they recommend using spatial transformations (such as contraction, flipping) for defense. Compared to previous methods, this method is more efficient because it requires almost no additional computational cost.

4.2. Defense Against Model Poisoning

(1) Filtering: Similar to defense methods against poisoned data, defense methods against poisoned models also begin

with model filtering. Fung et al. propose FoolsGold [25], which checks for and eliminates suspicious updates during local updates. FoolsGold is based on the fact that when a global model is trained by a group of attackers, they will submit updates with the same backdoor objectives throughout the training process, resulting in similar behaviors.




Figure 9. In the first stage, activation values of samples in each class are clustered into two groups. And in the second stage, it is determined which clusters correspond to poisoned samples.

However, this similarity does not occur among honest participants because each user's training datasets are unique and not shared with others. Therefore, malicious attackers can be separated from benign attackers through gradient updates. After detecting such anomalies, FoolGold maintains the learning rate of benign users (submitting only unique gradient updates) and reduces the learning rate of malicious users (repeatedly uploading similar gradient updates) to mitigate backdoor attacks. However, experimental results show that FoolsGold cannot defend against adaptive

attacks.




Figure 10. Februus uses a GAN-based repair method to reconstruct the masked regions to mitigate their adverse effects (such as benign accuracy reduction).

Li et al. [43] propose a spectral anomaly detection framework for a central aggregator that detects and erases malicious updates through strong detection model detection. The key idea of spectral anomaly detection is that there is a significant difference between the embedding of benign updates and backdoor updates in a low-dimensional latent space. One practical method for approximating low-dimensional embedding is to build a model using an encoder-decoder structure, where the encoder takes in the raw update and returns a low dimensional embedding, and the decoder is fed the embedding and outputs the generation error. After training the encoder-decoder model on benign updates, it can be used to identify backdoor updates, generating errors much higher than benign errors; malicious updates will be excluded from the aggregation process. However, this defense method cannot handle multi-trigger backdoor attacks, i.e., injecting various backdoors simultaneously.

Nguyen et al. propose FlGuard [60], a two-layer defense method that checks for locally updated updates with clear backdoor effects and eliminates residual backdoors through pruning, smoothing, and noise addition. Unlike FoolsGold [25], it also applies to multi-trigger backdoor attacks while maintaining high prediction accuracy for benign primary tasks. Additionally, FLDetector [104] proposes a method to detect malicious clients by checking the consistency of model updates. Essentially, the server predicts the client's model update based on past updates in each iteration. If the model update received by the client differs significantly from the predicted update over multiple iterations, the client is marked as malicious. Overall, the focus of these methods is mainly divided into two types: the first is to remove toxic updates from malicious clients before model aggregation, and the second is to reduce the impact of malicious clients on the aggregated model, such as reducing the

learning rate of suspicious clients.

(2) Robust Training: After filtering techniques, another class of techniques aims to directly mitigate backdoor attacks during model training through robust joint training. Differential privacy algorithms have been shown to be effective against backdoors [59], but they may compromise model performance under data imbalance commonly found in federated learning [4]. DP-FedAvg [55] (Central-DP) is a differential private aggregation strategy that eliminates outliers by clipping the norm of model updates and adding Gaussian noise, but the required amount of noise significantly reduces task accuracy. Sun et al. [66] proposed weak DP, which adds sufficient Gaussian noise to defeat backdoors while maintaining task accuracy, but it is ineffective against constraint-based backdoor attacks [44]. Additionally, DP-based defenses can potentially impact the benign performance of the global model, as the clipping factor also changes the weights of benign model updates.

In addition to DP-based defenses, Andreina et al. [3] propose Feedback-based Federated Learning (BaFFle) to eliminate backdoors. The key idea of BaFFle is to leverage participants to verify the global model. BaFFle includes a super-digit verification process for each round of federated learning. Specifically, each selected participant checks the current global model by computing a verification function on their secret data and reports to the central server whether the model is backdoored or not. The central server then decides whether to accept or reject the current global model based on feedback from all users. The verification function compares the error rate of a specific class of the current global model with the error rate of previously accepted global models. If the error rate is significantly different, the central server rejects the current global model as it may be backdoored and issues an alert. Unlike anomaly detection, BaFFle is compliant with secure aggregation.

Considering that all of the above defense works lack robustness certification, Xie et al. [92] propose the first general defense framework CRFL for training certifiably robust FL models against backdoor attacks. As shown in Fig. 11, CRFL controls the model smoothness through pruning and smoothing of model parameters and generates sample robustness certification against amplitude-limited backdoor attacks. Smoothness and perturbation methods are also used as additional components to constrain the L2 norm of individual updates to enhance defense performance [61]. In addition, the FL-WBC [75] method aims to identify the fragile parameter space in FL and perturb it during client training. FL-WBC also provides robustness guarantees against backdoor attacks and convergence guarantees for FedAvg. These developments demonstrate promising steps toward improving FL's robustness against backdoor attacks. In FLARE [85], a trust evaluation method is proposed that calculates a trust score for each model update based on the




Figure 11. Overview of certifiably robust federated learning (CRFL). CRFL controls the model smoothness through pruning and smoothing of model parameters and generates sample robustness certification against amplitude-limited backdoor attacks. Smoothness and perturbation methods are also used as additional components to constrain the L2 norm of individual updates to enhance defense performance.

difference between all model updates and their penultimate layer representation values. FLARE assumes that most clients are trustworthy and assigns low scores to updates that are far from benign update clusters. The model updates are then aggregated with their trust scores as weights and the global model is updated accordingly. In [46], the concept of anti-backdoor learning is introduced, which involves training a clean model given infected data, dividing the overall learning task into a dual task of learning the clean part of the data and the backdoor part. The article leverages two inherent weaknesses of backdoor attacks: models learn backdoor data faster than clean data, and the stronger the attack, the faster the model converges on the backdoor data. Additionally, backdoor tasks are mutually associated with specific classes. Based on these two weaknesses, a generic learning scheme is proposed to automatically prevent backdoor attacks during training. A two-stage gradient ascent mechanism is introduced to isolate and separate backdoor samples from the target class during the early training stage and break the association between backdoor samples and the target class during the later training stage.

(3) Model Reconstruction: The model reconstruction-based approach aims to eliminate hidden backdoors in infected models by directly modifying suspicious models. Therefore, even if triggers are included in the attack samples, the reconstructed model will still correctly predict them because the hidden backdoors have been removed. As mentioned earlier, the forgetting mechanism of federated backdoor attacks means that as training and model aggre-

gation progress, the backdoor will be forgotten in successive iterations. As a defense, this forgetting mechanism can also be utilized to create many defensive methods. Zeng et al. [97] define multiple training as a min-max problem and used implicit hyper-gradients to explain the interdependence between internal and external optimization. In [106], Zhao et al. show that hidden backdoors infecting DNNs can be repaired using pattern connectivity techniques with a certain number of benign samples. Li et al. [96] perturb backdoor-related neurons based on the distillation process, reconstructed (infected) DNNs using knowledge distillation techniques, and thus removed hidden backdoors. Huang et al. [37] propose a distillation technique that utilizes Cognitive Distills to extract Cognitive patterns. This is because the patterns of backdoor examples are generally small and sparse, making it possible to detect poisoned examples.

In addition to directly eliminating hidden backdoors, defense based on trigger synthesis first synthesizes backdoor triggers, and then in the second phase, eliminates hidden backdoors by suppressing the influence of triggers. These defenses have some similarities in the second phase with reconstruction-based defenses. For example, pruning and retraining are commonly used techniques to remove hidden backdoors in both defenses. However, compared to reconstruction-based defenses, the trigger-based defenses' trigger information makes the removal process more effective and efficient.

A GAN-based method is proposed in [111] to synthesize trigger distributions. In [34], they show that the detection process used in [82] to determine synthesized triggers has several failure modes and proposed a new defense method based on this observation. Additionally, Cheng et al. [93] reveal that the ∞ norm of activation values can be used to distinguish backdoor-related neurons based on synthesized triggers. Therefore, they propose performing an ∞ -based neuron pruning to remove neurons with high activation values in response to triggers. Similarly, Aiken et al. [2] also propose removing hidden backdoors by pruning DNNs based on synthesized triggers from another perspective.

In general, we do not consider filtering techniques to be a viable solution for mitigating backdoor attacks. Because filtering technologies are generally designed for specific types of attacks, they can be easily spoofed by attackers. Instead, we believe that the focus of research on backdoor defense should be on model reconstruction and robust training.

5. Byzantine Attack

The Byzantine attack is a malicious attack that typically occurs in distributed systems with malicious participants. In the field of computer science and cryptography, the Byzantine Generals Problem describes the scenario of distributed systems with malicious participants. The essence of the

Byzantine Generals Problem lies in how to achieve consensus among all nodes in the system when facing potential malicious participants. Byzantine attacks pose a significant threat to the security and reliability of distributed systems. This chapter focuses on categorizing Byzantine attacks into three types and provides an overview of the challenges faced and future directions for development. As shown in Fig. 13, we will analyze Byzantine attack from 3 aspects.

5.1. Sybil Attack

The first type of Byzantine attacks is the Sybil attack. It was first proposed by John Douceur et al. in 2002 [20]. As shown in Fig. 14, it is a type of network security attack where the attacker creates multiple fake identities or nodes to undermine the network's trust mechanism and disrupt its normal operation. The principle behind this attack is that the attacker creates a large number of fake identities, nodes, or accounts that appear independent and genuine to the network but are actually controlled by the attacker. The attacker can use various means to create these fake entities, including using forged IP addresses, anonymous proxy servers, virtual machines, and more. This attack mainly targets systems that rely on trust and identity verification mechanisms, such as peer-to-peer networks, social networks, and blockchains. Attackers can use a large number of fake identities to control the resources of the system, deceive other users, and disrupt consensus mechanisms. For example, in a peer-to-peer network, attackers can create numerous fake nodes to control the distribution process of files, leading to uneven resource distribution or network performance degradation.

The Sybil attack has posed significant threats to various network security systems since it was proposed. Bhagoji et al. [9], Jun et al. [49] and Yang et al. [94] explore it from the aspects of attack strategy, attack form and attack field respectively. Bhagoji et al. [9] explore some attack strategies against deep neural networks by enhancing malicious agent updates and employing alternating minimization strategies for stealth and adversarial targets. They demonstrate the possibility of effective and covert model poisoning attacks. Jun et al. [49] are the first to explore the Freerider attack in federated learning and propose an incremental weight attack method that can evade most defense monitoring methods at that time. They also introduce a new high-dimensional anomaly detection method called STD-DAGMM, which is particularly useful for detecting freeriders. This method has potential applications for detecting other models' weight anomalies as well, which will be discussed in the next chapter. Yang et al. [94] focus on the problem of vehicular ad hoc networks and criticized previous studies for only considering the security threats and impacts of Sybil attacks in VANETs without analyzing the po-

tential threats and impacts in VFCs (Vehicular Fog Computing). Therefore, the authors summarize four types of Sybil attacks that could affect VFCs in areas such as routing, vehicle decision-making, voting, and reputation systems.

In recent years, Sybil attacks have undergone significant development. Attackers employ more advanced techniques to create false identities, making Sybil nodes more realistic and difficult to distinguish. For instance, they may use virtualization technology to create multiple independent virtual machines, each with its own IP address and network identifier. Attackers can also utilize social engineering tactics and information acquisition methods to disguise themselves as genuine users or obtain additional fake identities through the authorization of legitimate users. Furthermore, attackers might construct more complex and realistic Sybil networks by leveraging shared information in social networks. With the widespread use of mobile devices and the wireless networks, Sybil attacks have also started to increase in these environments. Attackers can threaten the security of location services, social networks, and mobile applications by utilizing false identities and location information on mobile devices. In the future, Sybil attacks can be more extensively applied to blockchain technology, which enable attackers to manipulate consensus algorithms, launch double-spending attacks, or undermine the fairness of blockchain networks by controlling a large number of counterfeit identities.

5.2. Transmission Process Attack

The second type of Byzantine attacks is based on the transmission process. Byzantine attacks can occur during the communication process between participants. Attackers can tamper with, delete, or insert malicious messages to disrupt the transmission of model updates.

Zhang et al. [102] propose a centralized related probability small-scale attack (CDPS), which utilizes additional information. For instance, fusion rules in fault diagnosis [39] can be considered as existing knowledge usable for optimizing attack tactics. One form of CDPS attack involves collaborative efforts among malicious users who engage in communication. Moreover, sharing of information can assist malicious users in devising efficient attack strategies. Generally, colluding malicious users [23, 67] first exchange measurement values to ensure more accurate decisions on the licensed channel state and then report the coordinated forged results to enhance attack power. Moreover, Fatemieh et al. [24] introduce another attack model called the full-knowledge attack.

Cooperative spectrum sensing (CSS) is considered a promising method for identifying available spectrum. However, it not only requires a significant amount of communication resources but also introduces vulnerabilities to Byzantine attacks. Wu et al. [88] mention that attackers keep reporting false information to the fusion center consis-




Figure 12. The schematic of defense methods of Byzantine attack.




Figure 13. The schematic of Byzantine attack.




Figure 14. The schematic of the Sybil attack. Malicious users create a large number of fake nodes and control them to disrupt the normal operation of the network.

tently. In fact, attackers may attack discontinuously while appearing normal during rest periods. The authors propose a low-complexity sequential 0/1 (S0/1) method for CSS in the presence of strategic Byzantine attacks. It does not require strong assumptions or any prior knowledge, and this method will be discussed in the next chapter.

With the development of IoT devices, attackers may exploit communication vulnerabilities and weaknesses between IoT devices to launch future attacks. Attackers may also utilize high-speed, low-latency 5G/6G networks to carry out large-scale distributed denial-of-service (DDoS) attacks, bypass security measures using network slicing and virtualization technologies, or exploit vulnerabilities in mobile communication protocols to attack devices and sys-

tems.

5.3. Model Aggregation Attack

The third type of Byzantine attacks is known as model aggregation attack. In this type of attack, Byzantine attackers manipulate the model parameters of the participants in order to influence the aggregation results.




Figure 15. The schematic of a DDoS attack. Through the central control system, the attacker uses specific software or malicious code to remotely control multiple zombies, making them send a large number of requests to the target system at the same time to exhaust the processing power, network bandwidth or resources of the target system, thus causing the target system to fail to work normally.

As discussed by El Mhamdi et al. [31], these attacks by Byzantine workers may not prevent the model from converging but can instead lead to convergence on suboptimal solutions. This is a new way of thinking and there is a lot

of room for development in the future. However, most attacks disrupt the convergence of the model. Xie et al. [89] propose a form of attack called Gamber. In this attack, an attacker can modify some of the data communicated between participants. The attacker randomly selects data and makes malicious changes to them. The authors also mention several other attacks, such as omniscient. The attacker needs to have knowledge of the gradients sent by all the staff members and replaces some gradient vectors by taking the sum of all gradients, scaled by large negative values. The objective is to mislead the Stochastic Gradient Descent (SGD) algorithm into moving in the opposite direction with larger step sizes. Weaker attacks, such as Gaussian attacks, are also possible. In these attacks, some gradient vectors are replaced with random vectors sampled from a Gaussian distribution that has a large variance. Such attackers do not require any information from the staff members.




Figure 16. The schematic of Gaussian attack and Omnipotent attack. In a Gaussian attack, an attacker enters incorrect packets into the network, affecting the normal running of network flows. An omniscient attack is a hypothetical scenario where an adversary has complete and perfect knowledge or awareness of the system they are targeting. In this type of attack, the adversary has access to all the information and capabilities necessary to exploit vulnerabilities in the system. With this all-encompassing knowledge, the attacker can strategically plan and execute attacks to maximally exploit the system's weaknesses.

In many applications, multiple sources can provide descriptions of the same object or event, which can inevitably lead to conflicts in data or information. Resolving conflicts in data is an important challenge in achieving convergence in federated learning. Byzantine attackers have the ability to exploit these issues in order to manipulate the convergence of the model. Li et al. [42], Hsu et al. [36], and Jiang et al. [38] have all explored questions related to data differences and conflicts.

In future attacks, attackers may attempt to achieve their objectives by making modifications to the trained model. This can involve incorporating backdoors into the model, which would produce incorrect outputs under specific conditions, or tampering with the model parameters to disrupt its robustness. In order to evade detection, future attack strategies could involve the development of more advanced and covert techniques to camouflage malicious modifications. Additionally, malicious insiders who have access to the model parameters may exploit this privilege to launch attacks. They could manipulate the parameters to compromise the model's performance or even insert sensitive information within the model itself. In the coming years, attackers may explore more complicated techniques and algorithms to bypass security monitoring and detection mechanisms. The aim would be to implement internal personnel attacks more efficiently. These advancements would allow attackers to bypass existing safeguards and carry out their malicious activities with greater precision and effectiveness.

6. Defenses Against Byzantine Attack

With the advancement of Byzantine attacks, the corresponding defense methods are constantly evolving and being updated. As shown in Fig. 12, based on the different types of Byzantine attacks discussed in the previous chapter, we categorize the Byzantine defense methods into three classes:

6.1. Defenses against Sybil Attack

For Sybil attacks, the defense methods include trust-based evaluation methods and machine learning models for detecting fake nodes. Gil et al. [27] propose a novel algorithm that analyzes received wireless signals to detect the presence of deceptive clients created by adversaries. It does not require specialized hardware or private key exchanges, commercial Wi-Fi cards and software are enough. It utilizes the physical characteristics of wireless signals to "perceive" deceivers. The authors conduct experiments using the AscTec quadcopter server team and iRobot Create ground clients, which show a deception detection rate exceeding 96%. Blanchard et al. [10] introduce Krum, which detects and removes outliers in gradient aggregation, allowing convergence even in the presence of multiple Byzantine attackers, and with low complexity. Gil et al. [27] and Blanchard et al. [10] reduce the impact of the attacker, but don't accurately detect the attacker and remove it.

Both of the following methods detect the malicious attacker and remove its effects. Muñoz-González et al. [58] propose a new algorithm for robust federated learning called Adaptive Federated Averaging, aims to detect faults, attacks, and malicious updates provided by participants in a collaborative model. The authors also propose a Hidden Markov model to simulate and learn the quality of model updates provided by each participant during training. In

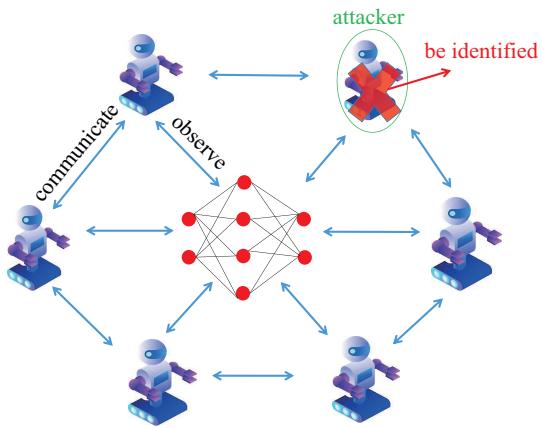


Figure 17. Robots can communicate with neighbors to get information, combined with their own information, in a certain round can eliminate the interference of malicious robots, reach a consensus.

the work of Mallmann-Trenn et al. [53], a combination of wireless signal analysis and observations from sociological learning is used to demonstrate rapid convergence of correct trust values for all robots in the team when facing attacks. All robots develop their own opinions about network trust by observing messages sent through the network as shown in Fig. 17. They compare their opinions with those of their neighbors to reach a consensus on whether the robots can be trusted. By utilizing their neighbors' opinions, each robot effectively increases the number of observations available to them about the network, thereby eliminating messages with a high probability of coming from malicious robots through cross-validation. Experiments show that in a limited number of communication rounds, all robots agree on the global consensus of trustworthiness in their neighborhood.

Sybil attacks pose a significant challenge as attackers can create a large number of false identities, making it difficult to detect and differentiate them. Therefore, developing secure systems requires consideration of multiple defense methods and selecting appropriate solutions based on specific application contexts. The future outlook for such defense techniques primarily focuses on the following areas: enhancing identity verification mechanisms to distinguish genuine users from Sybil nodes and establishing stronger trust mechanisms to identify and filter out Sybil nodes displaying malicious behavior. In recent years, the emergence of blockchain technology has provided a decentralized, tamper-resistant, and secure shared ledger, offering a potential defense against Sybil attacks. By combining identity verification, trust mechanisms, and blockchain technology, the ability to detect and prevent Sybil nodes can be improved.

6.2. Defense Against Transmission Process Attack

Defense methods against attacks during the transmission process include designing secure communication protocols and encryption mechanisms. Previous malicious detection and suppression algorithms focused only on simple "always attack" scenarios, where attackers continuously report false information to the central entity. In reality, attackers may intermittently attack while behaving normally at other times. Under the assumption of simple attack strategies, Chen et al. [15] propose sequential hypothesis testing to defend against Byzantine attacks, but it requires significant computation and prior knowledge. Therefore, Wu et al. [88] introduce S0/1 using support vector regression (SVR) to offset strategic Byzantine attack defense. Even in a blind scenario, this approach greatly reduces the sample size while providing higher correct sensing ratios. MuñozGonzález et al. [58] also propose a robust aggregation rule that detects and discards malicious or poor local model updates during each training iteration. It includes mechanisms for blocking unwanted participants, thereby enhancing computational and communication efficiency.

In traditional federated learning, the performance of models may decline due to variations in data distributions across devices. To solve this problem, Mostafa et al. [57] propose a technique called "representation matching." This technique improves model performance by matching features between the global model and local models. For clients with homogeneous data distributions, this approach consistently improves accuracy. For heterogeneous clients, in addition to improving accuracy, this approach enhances training robustness and avoids catastrophic training failures without the need for manual hyperparameter tuning for each task.




Figure 18. DDOS defense method of secret moving proxy nodes. The proxy node is constantly moving its position, confusing the attacker so that it cannot find the real target server.

Wang et al. [83] propose a mobile target defense mechanism that specifically addresses DDoS attacks targeting authenticated clients of Internet services. This mechanism utilizes a set of dynamic and hidden proxies to relay traffic between authenticated clients and servers. As shown in Fig. 18, by continuously replacing the attacked proxies with backup proxies and reallocating attacked clients to new proxies, innocent clients are separated from malicious insiders through a series of shuffling. In recent years, significant efforts have been made to defend against lowrate DDoS attacks. Attackers easily launch complex lowrate DDoS attacks by exploiting prominent features of cloud computing. Therefore, research on various DDoS attacks and their corresponding defense methods is crucial in protecting cloud infrastructure from devastating impacts of DDoS attacks. Agrawal et al. [1] conduct a comprehensive classification of all possible variants of cloud DDoS attack solutions and provide detailed insights into characterization, prevention, detection, and mitigation mechanisms.

In the future, these defense methods can be further developed by constructing more secure and robust encryption protocols and algorithms. Additionally, more effective Identity verification technologies can be developed to ensure the trustworthiness of communication parties. Continual improvement and enhancement of transport layer protocols, such as SSL/TLS, can provide stronger security and defense capabilities, including the ability to resist man-in-the-middle attacks, data tampering, and replay attacks. The development of real-time threat detection systems capable of detecting and responding to attacks during the transmission process can also be explored. With the advancement of quantum computing, traditional encryption algorithms may be threatened. Therefore, researching and developing quantum-safe communication protocols and algorithms is also important to ensure security during the transmission process and protect data confidentiality in a quantum computing environment.

6.3. Defense Against Model Aggregation Attack

Defense methods against model aggregation attacks include designing robust aggregation algorithms and using differential privacy techniques to protect model parameters.

El Mhamdi et al. [31] argue that mere convergence of the Byzantine resilient rules is not enough. It is possible that Byzantine workers' attacks can lead to convergence to the worst suboptimal solution. To address this, they propose a generic enhancement method called Bulyan. This method significantly reduces the leeway for Byzantine workers, constraining them within narrow boundaries. For common batch sizes, Bulyan achieves performance comparable to average speed.

Li et al. [42], Hsu et al. [36] and Jiang et al. [38] work on issues about data heterogeneity and conflicts. Li et al. [42]

model the problem using an optimization framework where ground truth and source reliability are defined as two sets of unknown variables. The objective is to minimize the overall weighted differences between the ground truth and multiple-source observations, with each source weighted by its reliability. This framework can incorporate different loss functions to identify features of various data types and has developed efficient computing methods that effectively identify the true information among conflicting data sources. Hsu et al. [36] present a method for synthesizing datasets with continuous identical ranges and provide performance metrics for the federated averaging algorithm. Experimental results demonstrate that performance deteriorates with distribution differences. The proposed method suppresses oscillations through the accumulation of gradient history, and it has been shown that using momentum on top of SGD for non-iid problems has achieved significant success in accelerating network training. Jiang et al. [38] introduce the combination of handling imbalanced data and Byzantine attacks in the context of federated learning for the first time. The concept of an intelligence pool is proposed to separate the task of judging the shared local model values from the aggregator, avoiding the limitations of a single criterion through two-layer verification shown in Fig. 19.

Xie et al. [89], Portnoy et al. [65] and Pillutla et al. [63] propose several rules or algorithms to make the model aggregation work smoothly. Xie et al. [89] introduce three robust aggregation rules for stochastic gradient descent (SGD) with low computational costs. These rules are the first to be theoretically and empirically studied under the non-convex setting, based on median aggregation. Portnoy et al. [65] propose a novel method called Byzantine Robust Client Weighting (BRCW) to minimize the impact of malicious clients in federated learning. BRCW assigns different weights to each client based on their credibility and contribution to the overall model accuracy. During training, the weights are dynamically updated according to each client's performance and behavior. Unlike previous work, this study considers the sample size provided by clients as untrustworthy, as it may come from malicious clients. BRCW exhibits robustness against various forms of attacks, including communication disruptions and data poisoning. Pillutla et al. [63] design a novel robust aggregation oracle based on classical geometric medians and demonstrate its robustness in federated learning with limited heterogeneity, even when up to half of the devices are faulty. The proposed RFA algorithm outperforms the standard FedAvg (14) in high damage scenarios and nearly matches FedAvg's performance in low damage scenarios, but with 1-3 times higher communication costs.

Future prospects for such defense methods will focus on more stricter authentication and authorization mechanisms, as well as parameter integrity protection. Mechanisms such




Figure 19. The structure of two-layer aggregation method. Participants share the global model, and models discarded by the filter layer are not directly abandoned, but are voted by participants to participate in the final aggregation process.

as Hash algorithms or digital signatures can be adopted to safeguard the integrity of parameters, preventing tampering during transmission and storage. Detection of malicious parameter modifications and corresponding response measures are also required. By comprehensively applying these methods, the defense capability against malicious parameter modification attacks can be enhanced, ensuring the security and integrity of systems and applications.

7. Adversarial Attack

Both adversarial attacks and backdoor attacks are techniques used to modify benign testing samples in order to make models misbehave during the inference process. While adversarial perturbations are sample-agnostic in universal adversarial attacks, these attacks can appear similar to backdoor attacks. Consequently, researchers who are unfamiliar with backdoor attacks may question their significance, as these attacks require additional controls on the training process to some extent. However, despite certain similarities, these attacks have essential differences.

Firstly, adversarial attackers need to control the inference process to a certain degree but not the training process of models. They must query the model results or even gradients multiple times to generate adversarial perturbations by optimization, given a fixed targeted model. On the other hand, backdoor attackers require modifying certain training stages, such as data collection and model training, without any additional requirements in the inference process. Secondly, from the perspective of attacked samples, backdoor attackers use known, non-optimized perturbations, whereas adversarial attackers require obtaining them through the op-

timization process based on the model output. This optimization in adversarial attacks requires multiple queries, making them unable to be real-time in many cases. Finally, the mechanisms of these attacks are fundamentally different. Adversarial vulnerability results from the differences in behaviors of models and humans, while backdoor attackers exploit the excessive learning ability of deep neural networks (DNNs) to establish a latent connection between trigger patterns and target labels.

In federated learning, because the data is not shared among participants, attackers can use this to generate adversarial samples locally, then inject them into the data sets of participants, and train the model with the data of other participants during the training process. So, how should attacks targeting federated learning be designed? In centralized learning, FGSM [29] is initially used to generate adversarial examples. Let θ be the parameters of a model, x the input to the model, η the perturbation to original input and $\|\eta\|_\infty \leq \epsilon$, y the targets associated with x (for machine learning tasks that have targets) and $J(\theta, x, y)$ be the cost used to train the neural network. We can linearize the cost function around the current value of θ , obtaining an optimal max-norm constrained perturbation of:

$$\eta = \epsilon sign(\nabla_x J(\theta, x, y)) \quad (1)$$

In federated learning, attackers can use FGSM locally to generate adversarial samples and then inject them into the datasets of some participants to affect the performance of the entire model.

In federated learning, attackers can use FGSM locally to generate adversarial samples and then inject them into the datasets of some participants to affect the performance of the entire model. PGD [52] is a more powerful adversarial sample generation method, which generates adversarial samples by iteratively perturbing the input. Specifically, PGD first randomly generates an initial disturbance δ_0 , and then iteratively updates the disturbance δ_t , to meet the constraint, that is, $|\delta_t|_\infty \leq \epsilon_{\max}$, where ϵ_{\max} is the maximum value of the disturbance. After each update, PGD also projects the disturbance δ_t back into the constraint space to ensure that it meets the constraint conditions. Specifically, the update formula for PGD is as follows:

$$x^{t+1} = Clip(x + sign(\nabla_x L(\theta, x', y)) \cdot \epsilon_{\max}, x - \epsilon_{\max}, x + \epsilon_{\max}), \quad (2)$$

Where Clip means to cut x into the interval, and ϵ_{\max} is the maximum value of the disturbance. Unlike FGSM, PGD requires multiple iterations to generate adversarial examples, so attackers need to use the local dataset and model parameters of each participant to iteratively update the perturbation.

In federated learning, using PGD (Projected Gradient Descent) or FGSM (Fast Gradient Sign Method) to generate

adversarial examples and perform adversarial attacks is a common method. However, compared to centralized learning, there are some considerations and techniques that need to be taken into account.

(1) Data Privacy and Distribution Heterogeneity: In federated learning, data is typically distributed across various edge devices, making it slightly more difficult to generate adversarial examples and perform attacks. Therefore, attackers need more complex strategies to adapt to the heterogeneity of the data distribution. In this case, the generation of adversarial examples should consider the data distribution and usage on each device. In federated learning, the amount and quality of data from each participant may vary, leading to imbalanced training data. This may affect the effectiveness of adversarial attacks, as attackers may not be able to generate sufficiently accurate adversarial examples to attack all participants' models.

(2) Network Latency and Bandwidth Limitations: In federated learning, devices communicate through a network, which may result in network latency or bandwidth limitations. Therefore, when generating adversarial examples and performing attacks, these factors need to be taken into consideration. For example, it may be necessary to optimize the size and generation speed of adversarial examples to adapt to network limitations. And in federated learning, the training steps of each participant are asynchronous, which may lead to synchronization issues. For example, attackers may not be able to perform attacks on all participants' models because they may be at different training steps or states.

(3) Update of Attack Strategies: Due to the dynamic nature of federated learning, attackers need to constantly update their attack strategies to adapt to changes in the learning model and data distribution. This may require designing more complex attack strategies and using more efficient optimization algorithms.

(4) Attack Detection and Defense: In federated learning, adversarial attacks may be easier to detect due to the distribution and usage of data. Therefore, it is necessary to design adversarial examples that are more difficult to detect or use more complex attack strategies to avoid detection.

Overall, using PGD or FGSM to generate adversarial examples and perform adversarial attacks in federated learning require considering the characteristics and limitations of federated learning, including the heterogeneity of data distribution, network latency and bandwidth limitations, update of attack strategies, and attack detection and defense. This may require designing more complex and efficient attack strategies and adversarial example generation methods.

8. Defenses against Adversarial Attack

Various adversarial examples have been great threats to deep learning models. Naturally, many defense methods are proposed to enhance the robustness of deep learning mod-

els. As adversarial samples share similarities with the previously discussed backdoor samples, defense methods involving sample filtering can be employed. However, in this section, we will refrain from discussing filtering methods and instead focus on the defense method of adversarial training. Adversarial training is a defense method against adversarial attacks. Its basic idea is to incorporate adversarial examples into the training process, enabling the model to better withstand unknown adversarial attacks. Specifically, adversarial training combines the original training dataset with adversarial samples generated to target the model, and retrain the model accordingly. This way, the model encounters adversarial examples continuously during the learning process, thereby enhancing its robustness and generalization capability.

As shown in Fig. 20, since adversarial training was first applied to federated learning in 2020, there are problems in four aspects. From these four aspects, we describe the current state of development of federated adversarial training. FAT (Federated Adversarial Training) is a method proposed in [112] that combines federated learning and adversarial training to reduce evasion threats during inference while preserving data privacy during training. In the FAT protocol, each device trains a local model and generates adversarial samples, which are mixed with the original data for training. The models are then uploaded to the server for aggregation. This helps to improve the robustness of the model against adversarial attacks. The authors note that the protocol does not work out of the box and requires careful tuning of the optimization parameters to achieve good results. Additionally, the authors acknowledge that the experiments were conducted on idealized federated settings and that further research is needed to evaluate the effectiveness of FAT in more realistic scenarios.

So, there are also several challenges with implementing adversarial training in federated learning, such as heterogeneous clients caused by distribution, certified guarantees, privacy leakage, and slow convergence rate.

8.1. Heterogeneous Clients Caused by Distribution

The major challenge in federated learning is the heterogeneity of clients. This means that the distribution of training data may vary among different clients, and the computational resources available to each client may also differ. The paper [68] proposes a novel method called FedRobust to address the problem of distribution shifts in federated learning. FedRobust is a gradient descent ascent (GDA) algorithm that solves the minimax robust optimization problem and can be efficiently implemented in a federated setting. This paper shows that FedRobust, which alternates between the perturbation and parameter model variables, will converge to a stationary point in the minimax objective that satisfies the Polyak-Łojasiewicz (PL) [64] condition. This op-

Adversarial Training in Federated Learning



Figure 20. The development of federated adversarial training.

timization method can be used to address the issue of distribution shifts in federated learning, which can significantly impact the performance of the trained model.




Figure 21. Decision boundary of DBFAT-trained model.

Zhang et al. [101] find that the accuracy of the federated learning model with adversarial training decreases on clean data, especially when the data of different clients are non-independent and identically distributed, so they propose an algorithm based on decision boundary as shown in Fig. 21, which uses local reweighting and global regularization to improve the accuracy and robustness in federated adversarial training. Chen et al. [13] study the skewed labels problem in federated learning and proposes a CalFAT framework to calculate the logits of each class. By using this framework, the models between different clients eventually converge, enhancing the robustness and accuracy of the global model. Hong et al. [35] have studied the problems caused by clients with different computing resources. During federated learning, some clients can support adversarial training, but some clients with lack of computing resources can

not support it. This paper studies how to spread resistance robustness from resource-rich users who can support confrontation training to those users with poor resources. This paper shows that an effective robust propagation algorithm is proposed by using BN layer technology correctly.

8.2. Certified Guarantees

Certified guarantees are needed in adversarial training to ensure reliable protection and security against unknown attacks in adversarial environments. Certified guarantees refer to rigorous proofs and assurances of model performance. They provide protection bounds against specific types of attacks, indicating that the model can maintain a certain level of performance regardless of how carefully the attacker designs the attack, such as high accuracy or robustness.

By using certified guarantees, adversarial training can provide a degree of predictability and security. This means that in practical deployment, there is higher confidence in the model's performance and trustworthiness, rather than relying solely on trial-and-error evaluation and protection. It offers a more reliable and definitive defense mechanism for models to handle unknown or complex attacks. Chen et al. [12] propose incorporating randomized smoothing techniques into federated adversarial training to enable data-private distributed learning with certifiable robustness to test-time adversarial perturbations. The experiments conducted in the paper show that this approach can deliver models as robust as those trained by centralized training, and can enable provably-robust classifiers to 2-bounded ad-

versarial perturbations in a distributed setup.

From the perspective of variance-bias decomposition, Zhou et al. [109] decompose the loss function into a combination of bias and variance, generate adversarial samples on the server and return them to each client for adversarial training. Then, they used any model aggregation algorithm to improve the global model’s adversarial robustness.

8.3. Privacy Leakage

As mentioned earlier, privacy protection is crucial in federated learning. In [100], it is proved that the model of confrontation training is more prone to privacy disclosure than the model of normal training. Zhou et al. [108] argue that conventional federated learning frameworks are vulnerable to strong adversarial attacks, even if adversarial training using locally generated adversarial examples is performed on each client. To address this problem, the paper proposes a new framework called FedBVA (Federated Learning with Bias-Variance Analysis) that provides a tiny amount of bias-variance perturbed data from the central server to the clients through asymmetrical communication. This approach dramatically improves the robustness of the training model under various settings, without violating the clients’ privacy.

8.4. Slow Convergence Rate

As described in [101], conducting adversarial training in federated learning is a challenging task because the convergence speed of adversarial training is very slow. In [72], a method is proposed to solve this problem. Assuming the local iteration number of federated learning is E, this article believes that a larger E will increase the drift between models and affect the adversarial robustness, but will converge quickly. A smaller E will increase the adversarial robustness but decrease the convergence efficiency. Therefore, this article proposes a dynamic mechanism for adjusting E, and uses FedCurv instead of FedAvg as the global aggregation algorithm, and proposes a new algorithm called FedDynAT to simultaneously improve the convergence speed and adversarial robustness. As shown in Tab.3, we list some of the methods mentioned above for comparison. It can be seen from the table that the adversarial robustness in federated learning is still weak (less than 40%). Therefore, the study of adversarial training in federated learning is still a very urgent matter.

9. Advanced Research and Problems

9.1. Stealthiness

A good attack method should not be easily detected by the defense system. Backdoor attackers need to effectively hide backdoors to make them difficult to detect. In Byzantine attacks, how do Byzantine attackers launch attacks while evading detection is also to be solved. Adversarial at-

Table 3. Comparison of FAT

Paper	Accuracy	Robustness
MixFAT	53.35%	29.14%
FedRBN	47.80%	26.87%
CalFAT	64.69%	35.03%
DBFAT	52.16%	27.80%
FedBVA	83.8%	21%
FedPGD	46.96%	28.7%
FedAvg	80.5%	9.9%

tacks need to generate adversarial samples without being detected. The attack methods developed above enhance their stealthiness by exploiting long-tail distribution, federated learning distribution, and dynamic generation techniques. Long tail distributed attacks and distributed attacks will eventually be discovered by the larger and larger model capacity, so how to develop a hidden attack method that can effectively face the large capacity model is a major problem in the current attack method research.

9.2. Trade-Off

There are many Trade-off issues in federated learning. In federated learning, client data is distributed among multiple participants. The protection of privacy is an important reason why federated learning is so popular. So it is an important challenge to ensure data privacy while implementing defenses. Effective defense mechanisms, such as fault-tolerant algorithms and majority voting, are needed to counter Byzantine attacks. However, these mechanisms need to strike a balance with factors like privacy preservation and computational efficiency. Adversarial training can enhance model robustness but may compromise generalization capability. Balancing the trade-off between adversarial robustness and generalization is a problem that requires exploration and resolution. Recently, some works have studied the latent connection between adversarial attacks and backdoor attacks. For example, Weng et al. [86] empirically demonstrated that defending against adversarial attacks via adversarial training may increase the risks of backdoor attacks. Solving these various Trade-off problems is also the focus of future research in federated learning.

9.3. Convergence

In defense methods, in order to improve the robustness of the model, existing methods often improve the training process of the model, such as distillation or adversarial training. These methods do play a good defense effect, but also bring the model convergence problems. The distributive na-

ture of federated learning further aggravates the problem of difficult convergence, because the distributive nature may cause problems such as heterogeneous data and unbalanced computing resources. In the methods mentioned in this paper, some articles have improved these problems in defense methods, such as using smoothing to produce equivalent results, or using reweighting to improve decision boundaries to solve the impact of data heterogeneity. But in general, how to make the model converge quickly is still a problem that needs to be studied in the future.

10. Conclusion

In this survey, we systematically introduce state of the art threats in federated learning systems, which mainly include byzantine attacks, backdoor attacks and adversarial attacks. And we also detailed describe corresponding defensive policies. We also introduce the advantages and disadvantages of these methods and how they work, and sort out the relationship between them. Additionally, we discuss a number of open problems of current defense methods, hoping them could help researchers identify and solve issues more quickly in the area of robust federated learning.

References

- [1] Neha Agrawal and Shashikala Tapaswi. Defense mechanisms against ddos attacks in a cloud computing environment: State-of-the-art and research challenges. *IEEE Communications Surveys & Tutorials*, 21(4):3769–3795, 2019. 14
- [2] William Aiken, Hyoungshick Kim, Simon Woo, and Jungwoo Ryoo. Neural network laundering: Removing black-box backdoor watermarks from deep neural networks. *Computers & Security*, 106:102277, 2021. 9
- [3] Sébastien Andreina, Giorgia Azzurra Marson, Helen Möllering, and Ghassan Karame. Baffle: Backdoor detection via feedback-based federated learning. In *2021 IEEE 41st International Conference on Distributed Computing Systems (ICDCS)*, pages 852–863. IEEE, 2021. 8
- [4] Eugene Bagdasaryan, Omid Poursaeed, and Vitaly Shmatikov. Differential privacy has disparate impact on model accuracy. *Advances in neural information processing systems*, 32, 2019. 8
- [5] Eugene Bagdasaryan, Andreas Veit, Yiqing Hua, Deborah Estrin, and Vitaly Shmatikov. How to backdoor federated learning. In *International conference on artificial intelligence and statistics*, pages 2938–2948. PMLR, 2020. 2, 3, 4, 5
- [6] Marco Barreno, Blaine Nelson, Russell Sears, Anthony D Joseph, and J Doug Tygar. Can machine learning be secure? In *Proceedings of the 2006 ACM Symposium on Information, computer and communications security*, pages 16–25, 2006. 3
- [7] Daniel Becking, Heiner Kirchhoffer, Gerhard Tech, Paul Haase, Karsten Müller, Heiko Schwarz, and Wojciech Samek. Adaptive differential filters for fast and communication-efficient federated learning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 3367–3376, 2022. 2
- [8] Jeremy Bernstein, Jiawei Zhao, Kamyar Azizzadenesheli, and Anima Anandkumar. signsgd with majority vote is communication efficient and fault tolerant. *arXiv preprint arXiv:1810.05291*, 2018. 3
- [9] Arjun Nitin Bhagoji, Supriyo Chakraborty, Prateek Mittal, and Seraphin Calo. Analyzing federated learning through an adversarial lens. In *International Conference on Machine Learning*, pages 634–643. PMLR, 2019. 10
- [10] Peva Blanchard, El Mahdi El Mhamdi, Rachid Guerraoui, and Julien Stainer. Machine learning with adversaries: Byzantine tolerant gradient descent. *Advances in neural information processing systems*, 30, 2017. 12
- [11] Bryant Chen, Wilka Carvalho, Nathalie Baracaldo, Heiko Ludwig, Benjamin Edwards, Taesung Lee, Ian Molloy, and Biplav Srivastava. Detecting backdoor attacks on deep neural networks by activation clustering. *arXiv preprint arXiv:1811.03728*, 2018. 2, 6
- [12] Cheng Chen, Bhavya Kailkhura, Ryan Goldhahn, and Yi Zhou. Certifiably-robust federated adversarial learning via randomized smoothing. In *2021 IEEE 18th International Conference on Mobile Ad Hoc and Smart Systems (MASS)*, pages 173–179. IEEE, 2021. 17
- [13] Chen Chen, Yuchen Liu, Xingjun Ma, and Lingjuan Lyu. Calfat: Calibrated federated adversarial training with label skewness. *Advances in Neural Information Processing Systems*, 35:3569–3581, 2022. 3, 17
- [14] Mingqing Chen, Ananda Theertha Suresh, Rajiv Mathews, Adeline Wong, Cyril Allauzen, Françoise Beaufays, and Michael Riley. Federated learning of n-gram language models. *arXiv preprint arXiv:1910.03432*, 2019. 2
- [15] Ruiliang Chen, Jung-Min Jerry Park, and Kaigui Bian. Robustness against byzantine failures in distributed spectrum sensing. *Computer Communications*, 35(17):2115–2124, 2012. 13
- [16] Yanbo Dai and Songze Li. Chameleon: Adapting to peer images for planting durable backdoors in federated learning. *arXiv preprint arXiv:2304.12961*, 2023. 5
- [17] Georgios Damaskinos, El-Mahdi El-Mhamdi, Rachid Guerraoui, Arsany Guirguis, and Sébastien Rouault. Aggregathor: Byzantine machine learning via robust gradient aggregation. *Proceedings of Machine Learning and Systems*, 1:81–106, 2019. 3
- [18] Bao Gia Doan, Ehsan Abbasnejad, and Damith C Ranasinghe. Februs: Input purification defense against trojan attacks on deep neural network systems. In *Annual computer security applications conference*, pages 897–912, 2020. 2, 7
- [19] Keval Doshi and Yasin Yilmaz. Federated learning-based driver activity recognition for edge devices. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 3338–3346, 2022. 2
- [20] John R Douceur. The sybil attack. In *International workshop on peer-to-peer systems*, pages 251–260. Springer, 2002. 10

- [21] David Enthoven and Zaid Al-Ars. An overview of federated deep learning privacy attacks and defensive strategies. *Federated Learning Systems: Towards Next-Generation AI*, pages 173–196, 2021. 2
- [22] Minghong Fang, Xiaoyu Cao, Jinyuan Jia, and Neil Gong. Local model poisoning attacks to {Byzantine-Robust} federated learning. In *29th USENIX security symposium (USENIX Security 20)*, pages 1605–1622, 2020. 3
- [23] Omid Fatemeh, Ranveer Chandra, and Carl A Gunter. Secure collaborative sensing for crowd sourcing spectrum data in white space networks. In *2010 IEEE Symposium on New Frontiers in Dynamic Spectrum (DySPAN)*, pages 1–12. IEEE, 2010. 10
- [24] Omid Fatemeh, Ali Farhadi, Ranveer Chandra, and Carl A Gunter. Using classification to protect the integrity of spectrum measurements in white space networks. In *NDSS*, 2011. 10
- [25] Clement Fung, Chris JM Yoon, and Ivan Beschastnikh. Mitigating sybils in federated learning poisoning. *arXiv preprint arXiv:1808.04866*, 2018. 7, 8
- [26] Arturo Geigel. Neural network trojan. *Journal of Computer Security*, 21(2):191–232, 2013. 4
- [27] Stephanie Gil, Swarun Kumar, Mark Mazumder, Dina Katabi, and Daniela Rus. Guaranteeing spoof-resilient multi-robot networks. *Autonomous Robots*, 41:1383–1400, 2017. 12
- [28] Xueluan Gong, Yanjiao Chen, Qian Wang, and Weihan Kong. Backdoor attacks and defenses in federated learning: State-of-the-art, taxonomy, and future directions. *IEEE Wireless Communications*, 2022. 3, 4
- [29] Ian J Goodfellow, Jonathon Shlens, and Christian Szegedy. Explaining and harnessing adversarial examples. *arXiv preprint arXiv:1412.6572*, 2014. 15
- [30] Tianyu Gu, Brendan Dolan-Gavitt, and Siddharth Garg. Badnets: Identifying vulnerabilities in the machine learning model supply chain. *arXiv preprint arXiv:1708.06733*, 2017. 4
- [31] Rachid Guerraoui, Sébastien Rouault, et al. The hidden vulnerability of distributed learning in byzantium. In *International Conference on Machine Learning*, pages 3521–3530. PMLR, 2018. 11, 14
- [32] Shangwei Guo, Tianwei Zhang, Han Yu, Xiaofei Xie, Lei Ma, Tao Xiang, and Yang Liu. Byzantine-resilient decentralized stochastic gradient descent. *IEEE Transactions on Circuits and Systems for Video Technology*, 32(6):4096–4106, 2021. 3
- [33] Shangwei Guo, Xu Zhang, Fei Yang, Tianwei Zhang, Yan Gan, Tao Xiang, and Yang Liu. Robust and privacy-preserving collaborative learning: A comprehensive survey. *arXiv e-prints*, pages arXiv–2112, 2021. 2
- [34] Wenbo Guo, Lun Wang, Yan Xu, Xinyu Xing, Min Du, and Dawn Song. Towards inspecting and eliminating trojan backdoors in deep neural networks. In *2020 IEEE International Conference on Data Mining (ICDM)*, pages 162–171. IEEE, 2020. 9
- [35] Junyuan Hong, Haotao Wang, Zhangyang Wang, and Jiayu Zhou. Federated robustness propagation: Sharing robustness in heterogeneous federated learning. *arXiv preprint arXiv:2106.10196*, 2021. 17
- [36] Tzu-Ming Harry Hsu, Hang Qi, and Matthew Brown. Measuring the effects of non-identical data distribution for federated visual classification. *arXiv preprint arXiv:1909.06335*, 2019. 12, 14
- [37] Hanxun Huang, Xingjun Ma, Sarah Erfani, and James Bailey. Distilling cognitive backdoor patterns within an image. *arXiv preprint arXiv:2301.10908*, 2023. 9
- [38] Yanna Jiang, Baihe Ma, Xu Wang, Guangsheng Yu, Caijun Sun, Wei Ni, and Ren Ping Liu. A secure aggregation for federated learning on long-tailed data. *arXiv preprint arXiv:2307.08324*, 2023. 12, 14
- [39] Bhavya Kailkhura, Yunghsiang S Han, Swastik Brahma, and Pramod K Varshney. Distributed bayesian detection with byzantine data. *arXiv preprint arXiv:1307.3544*, 2013. 10
- [40] Alexander Kirillov, Eric Mintun, Nikhila Ravi, Hanzi Mao, Chloe Rolland, Laura Gustafson, Tete Xiao, Spencer Whitehead, Alexander C Berg, Wan-Yen Lo, et al. Segment anything. *arXiv preprint arXiv:2304.02643*, 2023. 1
- [41] Leslie Lamport, Robert Shostak, and Marshall Pease. The byzantine generals problem. In *Concurrency: the works of Leslie Lamport*, pages 203–226. 2019. 3
- [42] Qi Li, Yaliang Li, Jing Gao, Bo Zhao, Wei Fan, and Jiawei Han. Resolving conflicts in heterogeneous data by truth discovery and source reliability estimation. In *Proceedings of the 2014 ACM SIGMOD international conference on Management of data*, pages 1187–1198, 2014. 12, 14
- [43] Suyi Li, Yong Cheng, Wei Wang, Yang Liu, and Tianjian Chen. Learning to detect malicious clients for robust federated learning. *arXiv preprint arXiv:2002.00211*, 2020. 2, 8
- [44] Xiaoxiao Li, Zhao Song, and Jiaming Yang. Federated adversarial learning: A framework with convergence analysis. In *International Conference on Machine Learning*, pages 19932–19959. PMLR, 2023. 3, 8
- [45] Yiming Li, Tongqing Zhai, Baoyuan Wu, Yong Jiang, Zhifeng Li, and Shutao Xia. Rethinking the trigger of backdoor attack. *arXiv preprint arXiv:2004.04692*, 2020. 5
- [46] Yige Li, Xixiang Lyu, Nodens Koren, Lingjuan Lyu, Bo Li, and Xingjun Ma. Anti-backdoor learning: Training clean models on poisoned data. *Advances in Neural Information Processing Systems*, 34:14900–14912, 2021. 9
- [47] Yiming Li, Tongqing Zhai, Yong Jiang, Zhifeng Li, and Shu-Tao Xia. Backdoor attack in the physical world. *arXiv preprint arXiv:2104.02361*, 2021. 7
- [48] Bill Yuchen Lin, Chaoyang He, Zihang Zeng, Hulin Wang, Yufen Huang, Mahdi Soltanolkotabi, Xiang Ren, and Salman Avestimehr. Fednlp: A research platform for federated learning in natural language processing. *arXiv preprint arXiv:2104.08815*, 2021. 2
- [49] Jierui Lin, Min Du, and Jian Liu. Free-riders in federated learning: Attacks and defenses. *arXiv preprint arXiv:1911.12560*, 2019. 10
- [50] Rui Liu, Pengwei Xing, Zichao Deng, Anran Li, Cuntai Guan, and Han Yu. Federated graph neural networks:

- Overview, techniques and challenges. *arXiv preprint arXiv:2202.07256*, 2022. 2
- [51] Yuntao Liu, Yang Xie, and Ankur Srivastava. Neural trojans. In *2017 IEEE International Conference on Computer Design (ICCD)*, pages 45–48. IEEE, 2017. 2, 7
- [52] Aleksander Madry, Aleksandar Makelov, Ludwig Schmidt, Dimitris Tsipras, and Adrian Vladu. Towards deep learning models resistant to adversarial attacks. *arXiv preprint arXiv:1706.06083*, 2017. 15
- [53] Frederik Mallmann-Trenn, Matthew Caversi, and Stephanie Gil. Crowd vetting: Rejecting adversaries via collaboration with application to multirobot flocking. *IEEE Transactions on Robotics*, 38(1):5–24, 2021. 13
- [54] Brendan McMahan, Eider Moore, Daniel Ramage, Seth Hampson, and Blaise Aguera y Arcas. Communication-efficient learning of deep networks from decentralized data. In *Artificial intelligence and statistics*, pages 1273–1282. PMLR, 2017. 1
- [55] H Brendan McMahan, Daniel Ramage, Kunal Talwar, and Li Zhang. Learning differentially private recurrent language models. *arXiv preprint arXiv:1710.06963*, 2017. 8
- [56] Chenglin Miao, Qi Li, Houping Xiao, Wenjun Jiang, Mengdi Huai, and Lu Su. Towards data poisoning attacks in crowd sensing systems. In *Proceedings of the Eighteenth ACM International Symposium on Mobile Ad Hoc Networking and Computing*, pages 111–120, 2018. 3
- [57] Hesham Mostafa. Robust federated learning through representation matching and adaptive hyper-parameters. *arXiv preprint arXiv:1912.13075*, 2019. 13
- [58] Luis Muñoz-González, Kenneth T Co, and Emil C Lupu. Byzantine-robust federated machine learning through adaptive model averaging. *arXiv preprint arXiv:1909.05125*, 2019. 12, 13
- [59] Mohammad Naseri, Jamie Hayes, and Emiliano De Cristofaro. Toward robustness and privacy in federated learning: Experimenting with local and central differential privacy. *arXiv preprint arXiv:2009.03561*, 2020. 8
- [60] Thien Duc Nguyen, Phillip Rieger, Mohammad Hossein Yalame, Helen Möllering, Hossein Fereidooni, Samuel Marchal, Markus Miettinen, Azalia Mirhoseini, Ahmad-Reza Sadeghi, Thomas Schneider, et al. Flguard: Secure and private federated learning. *Cryptography and Security*, (Preprint), 2021. 8
- [61] Thien Duc Nguyen, Phillip Rieger, Roberta De Viti, Huili Chen, Björn B Brandenburg, Hossein Yalame, Helen Möllering, Hossein Fereidooni, Samuel Marchal, Markus Miettinen, et al. {FLAME}: Taming backdoors in federated learning. In *31st USENIX Security Symposium (USENIX Security 22)*, pages 1415–1432, 2022. 8
- [62] Mustafa Safa Ozdayi, Murat Kantarcioglu, and Yulia R Gel. Defending against backdoors in federated learning with robust learning rate. In *Proceedings of the AAAI Conference on Artificial Intelligence*, pages 9268–9276, 2021. 3
- [63] Krishna Pillutla, Sham M Kakade, and Zaid Harchaoui. Robust aggregation for federated learning. *IEEE Transactions on Signal Processing*, 70:1142–1154, 2022. 14
- [64] Boris Teodorovich Polyak et al. Gradient methods for minimizing functionals. *Zhurnal vychislitel'noi matematiki i matematicheskoi fiziki*, 3(4):643–653, 1963. 16
- [65] Amit Portnoy and Danny Hendler. Towards realistic byzantine-robust federated learning. 2020. 14
- [66] Saurav Prakash and Amir Salman Avestimehr. Mitigating byzantine attacks in federated learning. *arXiv preprint arXiv:2010.07541*, 2020. 3, 8
- [67] Zhengrui Qin, Qun Li, and George Hsieh. Defending against cooperative attacks in cooperative spectrum sensing. *IEEE Transactions on Wireless communications*, 12(6):2680–2687, 2013. 10
- [68] Amirhossein Reisizadeh, Farzan Farnia, Ramtin Pedarsani, and Ali Jadbabaie. Robust federated learning: The case of affine distribution shifts. *Advances in Neural Information Processing Systems*, 33:21554–21565, 2020. 16
- [69] Nuria Rodríguez-Barroso, Daniel Jiménez-López, M Victoria Luzón, Francisco Herrera, and Eugenio Martínez-Cámara. Survey on federated learning threats: Concepts, taxonomy on attacks and defences, experimental study and challenges. *Information Fusion*, 90:148–173, 2023. 2
- [70] Ahmed Salem, Rui Wen, Michael Backes, Shiqing Ma, and Yang Zhang. Dynamic backdoor attacks against machine learning models. In *2022 IEEE 7th European Symposium on Security and Privacy (EuroS&P)*, pages 703–718. IEEE, 2022. 4
- [71] Ali Shafahi, W Ronny Huang, Mahyar Najibi, Octavian Suciu, Christoph Studer, Tudor Dumitras, and Tom Goldstein. Poison frogs! targeted clean-label poisoning attacks on neural networks. *Advances in neural information processing systems*, 31, 2018. 5
- [72] Devansh Shah, Parijat Dube, Supriyo Chakraborty, and Ashish Verma. Adversarial training in communication constrained federated learning. *arXiv preprint arXiv:2103.01319*, 2021. 18
- [73] Shiqi Shen, Shruti Tople, and Prateek Saxena. Auror: Defending against poisoning attacks in collaborative deep learning systems. In *Proceedings of the 32nd Annual Conference on Computer Security Applications*, pages 508–519, 2016. 5
- [74] Jacob Steinhardt, Pang Wei W Koh, and Percy S Liang. Certified defenses for data poisoning attacks. *Advances in neural information processing systems*, 30, 2017. 5
- [75] Jingwei Sun, Ang Li, Louis DiValentin, Amin Hassanzadeh, Yiran Chen, and Hai Li. Fl-wbc: Enhancing robustness against model poisoning attacks in federated learning from a client perspective. *Advances in Neural Information Processing Systems*, 34:12613–12624, 2021. 8
- [76] Yuwei Sun, Hideya Ochiai, and Jun Sakuma. Semi-targeted model poisoning attack on federated learning via backward error analysis. In *2022 International Joint Conference on Neural Networks (IJCNN)*, pages 1–8. IEEE, 2022. 2, 6
- [77] Ziteng Sun, Peter Kairouz, Ananda Theertha Suresh, and H Brendan McMahan. Can you really backdoor federated learning? *arXiv preprint arXiv:1911.07963*, 2019. 3
- [78] Asadullah Tariq, Mohamed Adel Serhani, Farag Sallabi, Tariq Qayyum, Ezedin S Barka, and Khaled A Shuaib.

- Trustworthy federated learning: A survey. *arXiv preprint arXiv:2305.11537*, 2023. 2
- [79] Vale Tolpegin, Stacey Truex, Mehmet Emre Gursoy, and Ling Liu. Data poisoning attacks against federated learning systems. In *Computer Security–ESORICS 2020: 25th European Symposium on Research in Computer Security, ESORICS 2020, Guildford, UK, September 14–18, 2020, Proceedings, Part I 25*, pages 480–501. Springer, 2020. 5
- [80] Brandon Tran, Jerry Li, and Aleksander Madry. Spectral signatures in backdoor attacks. *Advances in neural information processing systems*, 31, 2018. 2, 6
- [81] Paul Voigt and Axel Von dem Bussche. The eu general data protection regulation (gdpr). *A Practical Guide, 1st Ed., Cham: Springer International Publishing*, 10(3152676): 10–5555, 2017. 1
- [82] Bolun Wang, Yuanshun Yao, Shawn Shan, Huiying Li, Bimal Viswanath, Haitao Zheng, and Ben Y Zhao. Neural cleanse: Identifying and mitigating backdoor attacks in neural networks. In *2019 IEEE Symposium on Security and Privacy (SP)*, pages 707–723. IEEE, 2019. 9
- [83] Huangxin Wang, Quan Jia, Dan Fleck, Walter Powell, Fei Li, and Angelos Stavrou. A moving target ddos defense mechanism. *Computer Communications*, 46:10–21, 2014. 14
- [84] Hongyi Wang, Kartik Sreenivasan, Shashank Rajput, Harit Vishwakarma, Saurabh Agarwal, Jy-yong Sohn, Kangwook Lee, and Dimitris Papailiopoulos. Attack of the tails: Yes, you really can backdoor federated learning. *Advances in Neural Information Processing Systems*, 33:16070–16084, 2020. 3, 5
- [85] Ning Wang, Yang Xiao, Yimin Chen, Yang Hu, Wenjing Lou, and Y Thomas Hou. Flare: defending federated learning against model poisoning attacks via latent space representations. In *Proceedings of the 2022 ACM on Asia Conference on Computer and Communications Security*, pages 946–958, 2022. 8
- [86] Cheng-Hsin Weng, Yan-Ting Lee, and Shan-Hung Brandon Wu. On the trade-off between adversarial and backdoor robustness. *Advances in Neural Information Processing Systems*, 33:11973–11983, 2020. 18
- [87] Chuhan Wu, Fangzhao Wu, Yang Cao, Yongfeng Huang, and Xing Xie. Fedgnn: Federated graph neural network for privacy-preserving recommendation. *Cornell University - arXiv, Cornell University - arXiv*, 2021. 2
- [88] Jun Wu, Yue Yu, Tiecheng Song, and Jing Hu. Sequential 0/1 for cooperative spectrum sensing in the presence of strategic byzantine attack. *IEEE Wireless Communications Letters*, 8(2):500–503, 2018. 10, 13
- [89] Cong Xie, Oluwasanmi Koyejo, and Indranil Gupta. Generalized byzantine-tolerant sgd. *arXiv preprint arXiv:1802.10116*, 2018. 12, 14
- [90] Chulin Xie, Keli Huang, Pin-Yu Chen, and Bo Li. Dba: Distributed backdoor attacks against federated learning. In *International conference on learning representations*, 2019. 4
- [91] Cong Xie, Oluwasanmi Koyejo, and Indranil Gupta. Fall of empires: Breaking byzantine-tolerant sgd by inner prod-
- uct manipulation. In *Uncertainty in Artificial Intelligence*, pages 261–270. PMLR, 2020. 3
- [92] Chulin Xie, Minghao Chen, Pin-Yu Chen, and Bo Li. Crfl: Certifiably robust federated learning against backdoor attacks. In *International Conference on Machine Learning*, pages 11372–11382. PMLR, 2021. 8
- [93] Kaidi Xu, Sijia Liu, Pin-Yu Chen, Pu Zhao, and Xue Lin. Defending against backdoor attack on deep neural networks. *arXiv preprint arXiv:2002.12162*, 2020. 9
- [94] Haonan Yang, Yongchao Zhong, Bo Yang, Yiyu Yang, Zifeng Xu, Longjuan Wang, and Yuqing Zhang. An overview of sybil attack detection mechanisms in vfc. In *2022 52nd Annual IEEE/IFIP International Conference on Dependable Systems and Networks Workshops (DSN-W)*, pages 117–122. IEEE, 2022. 10
- [95] Qiang Yang, Yang Liu, Yong Cheng, Yan Kang, Tianjian Chen, and Han Yu. Federated learning: synthesis lectures on artificial intelligence and machine learning. *vol*, 13:1–207, 2019. 2
- [96] Kota Yoshida and Takeshi Fujino. Disabling backdoor and identifying poison data by using knowledge distillation in backdoor attacks on deep neural networks. In *Proceedings of the 13th ACM Workshop on Artificial Intelligence and Security*, pages 117–127, 2020. 9
- [97] Yi Zeng, Si Chen, Won Park, Z Morley Mao, Ming Jin, and Ruoxi Jia. Adversarial unlearning of backdoors via implicit hypergradient. *arXiv preprint arXiv:2110.03735*, 2021. 9
- [98] Yi Zeng, Won Park, Z Morley Mao, and Ruoxi Jia. Rethinking the backdoor attacks’ triggers: A frequency perspective. In *Proceedings of the IEEE/CVF international conference on computer vision*, pages 16473–16481, 2021. 2, 6
- [99] Hengtong Zhang, Tianhang Zheng, Jing Gao, Chenglin Miao, Lu Su, Yaliang Li, and Kui Ren. Data poisoning attack against knowledge graph embedding. *arXiv preprint arXiv:1904.12052*, 2019. 3
- [100] Jingyang Zhang, Yiran Chen, and Hai Li. Privacy leakage of adversarial training models in federated learning systems. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 108–114, 2022. 18
- [101] Jie Zhang, Bo Li, Chen Chen, Lingjuan Lyu, Shuang Wu, Shouhong Ding, and Chao Wu. Delving into the adversarial robustness of federated learning. *arXiv preprint arXiv:2302.09479*, 2023. 3, 17, 18
- [102] Linyuan Zhang, Guoru Ding, Qihui Wu, Yulong Zou, Zhu Han, and Jinlong Wang. Byzantine attack and defense in cognitive radio networks: A survey. *IEEE Communications Surveys & Tutorials*, 17(3):1342–1363, 2015. 10
- [103] Yifei Zhang, Dun Zeng, Jinglong Luo, Zenglin Xu, and Irwin King. A survey of trustworthy federated learning with perspectives on security, robustness, and privacy. *arXiv preprint arXiv:2302.10637*, 2023. 2
- [104] Zaixi Zhang, Xiaoyu Cao, Jinyuan Jia, and Neil Zhenqiang Gong. Fldetector: Defending federated learning against model poisoning attacks via detecting malicious clients. In *Proceedings of the 28th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*, pages 2545–2555, 2022. 8

- [105] Zhengming Zhang, Ashwinee Panda, Linyue Song, Yao-qing Yang, Michael Mahoney, Prateek Mittal, Ramchandran Kannan, and Joseph Gonzalez. Neurotoxin: Durable backdoors in federated learning. In *International Conference on Machine Learning*, pages 26429–26446. PMLR, 2022. [2](#), [6](#)
- [106] Pu Zhao, Pin-Yu Chen, Payel Das, Karthikeyan Natesan Ramamurthy, and Xue Lin. Bridging mode connectivity in loss landscapes and adversarial robustness. *arXiv preprint arXiv:2005.00060*, 2020. [9](#)
- [107] Xingchen Zhou, Ming Xu, Yiming Wu, and Ning Zheng. Deep model poisoning attack on federated learning. *Future Internet*, 13(3):73, 2021. [2](#), [6](#)
- [108] Yao Zhou, Jun Wu, and Jingrui He. Adversarially robust federated learning for neural networks. 2020. [18](#)
- [109] Yao Zhou, Jun Wu, Haixun Wang, and Jingrui He. Adversarial robustness through bias variance decomposition: A new perspective for federated learning. In *Proceedings of the 31st ACM International Conference on Information & Knowledge Management*, pages 2753–2762, 2022. [18](#)
- [110] Chen Zhu, W Ronny Huang, Hengduo Li, Gavin Taylor, Christoph Studer, and Tom Goldstein. Transferable clean-label poisoning attacks on deep neural nets. In *International Conference on Machine Learning*, pages 7614–7623. PMLR, 2019. [5](#)
- [111] Liuwan Zhu, Rui Ning, Cong Wang, Chunsheng Xin, and Hongyi Wu. Gangsweep: Sweep out neural backdoors by gan. In *Proceedings of the 28th ACM International Conference on Multimedia*, pages 3173–3181, 2020. [9](#)
- [112] Giulio Zizzo, Ambrish Rawat, Mathieu Sinn, and Beat Buesser. Fat: Federated adversarial training. *arXiv preprint arXiv:2012.01791*, 2020. [3](#), [16](#)